

```
%tensorflow_version 1.x
```

```
☞ TensorFlow 1.x selected.
```

```
import pandas as pd
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt
```

```
np.random.seed(451)
```

```
from keras.datasets import cifar10
```

```
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
x_train = x_train / 255.0
```

```
x_test = x_test / 255.0
```

```
☞ Using TensorFlow backend.
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500096/170498071 [=====] - 2s 0us/step
```

```
#gray = 0.2989 * r + 0.5870 * g + 0.1140 * b
```

```
x_train_gray = np.dot(x_train[:,:,:,:3], [0.299, 0.587, 0.114])
```

```
x_test_gray = np.dot(x_test[:,:,:,:3], [0.299, 0.587, 0.114])
```

```
x_train_gray = x_train_gray.reshape(-1,32,32,1)
```

```
x_test_gray = x_test_gray.reshape(-1,32,32,1)
```

```
from keras.utils.np_utils import to_categorical
```

```
y_train_cat = to_categorical(y_train)
```

```
y_test_cat = to_categorical(y_test)
```


```
plt.imshow(x_train[1])
```

```
plt.show()
```

```
plt.imshow(x_train_gray[1,:,:,:0], cmap='gray')
```

```
plt.show()
```

```
☞
```



```
np.random.seed(451)
```

```
import datetime
```

```
from keras.layers import Flatten, Activation, Conv2D, MaxPool2D, AvgPool2D, Dense, Dropout
from keras.optimizers import Adam, SGD
from keras.models import Sequential
import keras.backend as K
from keras.regularizers import l1, l2
from keras.callbacks import EarlyStopping, ModelCheckpoint, TensorBoard, ReduceLROnPlateau
from keras.models import model_from_json, Model
```

```
def build_tower(input_layer, features_nr, shape, tower_nr,
                dropout=False, normalization=False, regularization="l2", dropout_ratio=0.2
                #3x3 kernel tower
                tower = Conv2D(features_nr, (1,1), padding='same', activation='relu',
                                kernel_regularizer=regularization, name='tower_%d_%dx%da'%(tower_nr,
                                tower = Conv2D(features_nr*2, shape, padding='same', activation='relu',
                                kernel_regularizer=regularization, name='tower_%d_%dx%db'%(tower_nr,
                #condidional dropout/normalization
                if dropout:
                    tower = Dropout(dropout_ratio, name='tower_%d_%dx%ddrop'%(tower_nr, shape[0], shap
                if normalization:
                    tower = BatchNormalization(name='tower_%d_%dx%dnorm'%(tower_nr, shape[0], shape[1]

                return tower
```

```
def build_simple_tower(input_layer, features_nr, shape, tower_nr,
                       dropout=False, normalization=False, regularization="l2", dropout_ratio=0.2
                       #3x3 kernel tower
                       tower = Conv2D(features_nr, shape, padding='same', activation='relu',
                                       kernel_regularizer=regularization,
                                       name='tower_simple_%d_%dx%db'%(tower_nr, shape[0], shape[1]))(input_lay
                       #condidional dropout/normalization
                       if dropout:
                           tower = Dropout(dropout_ratio, name='tower_%d_%dx%ddrop'%(tower_nr, shape[0], shap
                       if normalization:
                           tower = BatchNormalization(name='tower_%d_%dx%dnorm'%(tower_nr, shape[0], shape[1]

                       return tower
```

```

def build_tower_subsample(input_layer, features_nr, shape, tower_nr,
                          dropout=False, normalization=False, regularization='l2', dropout_ratio=0.5):
    tower = build_tower(input_layer, features_nr, shape, tower_nr,
                        dropout, normalization, regularization, dropout_ratio)
    pool = MaxPooling2D((2,2), padding='same', name='tower_%d_2x2subsample'%(tower_nr))(tower)

    return pool

def build_simple_tower_subsample(input_layer, features_nr, shape, tower_nr,
                                 dropout=False, normalization=False, regularization='l2', dropout_ratio=0.5):
    tower = build_simple_tower(input_layer, features_nr, shape, tower_nr,
                              dropout, normalization, regularization, dropout_ratio)
    pool = MaxPooling2D((2,2), padding='same', name='tower_%d_2x2subsample'%(tower_nr))(tower)

    return pool

def build_dense(input_layer, neurons_nr, dense_nr,
                dropout=False, normalization=False, regularization='l2', dropout_ratio=0.5):
    dense = Dense(neurons_nr, kernel_regularizer=regularization,
                  name='dense_%d_%d'%(dense_nr, neurons_nr))(input_layer)

    if dropout:
        dense = Dropout(dropout_ratio, name='dense_%d_%ddrop'%(dense_nr, neurons_nr))(dense)
    if normalization:
        dense = BatchNormalization(name='dense_%d_%dnorm'%(dense_nr, neurons_nr))(dense)

    return dense

def build_inception_module(input_layer, features_nr, module_nr,
                           dropout=False, normalization=False, regularization='l2', dropout_ratio=0.5):
    #feature_nr is an array we'll use to build our layers
    #data is in the form: [1x1, 3x3 reduce, 3x3, 5x5 reduce, 5x5, pool proj]

    inception_1x1 = Conv2D(features_nr[0],1,1,border_mode='same',activation='relu',name='i%d_1x1'%(module_nr))
    inception_3x3_reduce = Conv2D(features_nr[1],1,1,border_mode='same',activation='relu',name='i%d_3x3_reduce'%(module_nr))
    inception_3x3 = Conv2D(features_nr[2],3,3,border_mode='same',activation='relu',name='i%d_3x3'%(module_nr))
    inception_5x5_reduce = Conv2D(features_nr[3],1,1,border_mode='same',activation='relu',name='i%d_5x5_reduce'%(module_nr))
    inception_5x5 = Conv2D(features_nr[4],5,5,border_mode='same',activation='relu',name='i%d_5x5'%(module_nr))
    inception_pool = MaxPooling2D(pool_size=(3,3),strides=(1,1),border_mode='same',name='i%d_pool'%(module_nr))
    inception_pool_proj = Conv2D(features_nr[5],1,1,border_mode='same',activation='relu',name='i%d_pool_proj'%(module_nr))
    inception_output = concatenate([inception_1x1,inception_3x3,inception_5x5,inception_pool_proj],axis=-1)

    if dropout:
        inception_output = Dropout(dropout_ratio, name='inception_%d_/output_drop'%(module_nr))(inception_output)
    if normalization:
        inception_output = BatchNormalization(name='inception_%d_/output_norm'%(module_nr))(inception_output)

    pooled = MaxPooling2D((2,2), padding='same', name='inception_%d_2x2subsample'%(module_nr))(inception_output)

    return pooled

```

```

return pooled

i='cifar10-nrcrt7-'+datetime.datetime.now().strftime("%I:%M%p_%B-%d-%Y")

K.clear_session()

!mkdir -p models
!mkdir -p logs

a = EarlyStopping(monitor='val_loss', min_delta=0, patience=10, verbose=1, mode='auto')#wi
b = ModelCheckpoint(monitor='val_loss', filepath='./models/'+str(i)+'.hdf5', verbose=1, sa
c = TensorBoard(log_dir='./logs/'+str(i),
                 write_grads=True,
                 write_graph=True,
                 write_images=True,
                 batch_size=128)#saves a log file for tensorboard; remember to save differ

#we'll use this instead of decay
d = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=5, verbose=0, mode='auto',

callbacks=[a,b,c,d]

#-----model definition-----

use_norm = True
lr_rate = 0.001

input_img = Input(shape = (32, 32, 3), name='input')

#conv_1 = Conv2D(1, (1,1), padding='same', activation='relu',
                # kernel_regularizer = regularization, name='conv_64x64x1_inception_in')(ir

#hopefully this will learn a good internal representation of the image channels
#conv_1 = Conv2D(1, (1,1), padding='same', activation='relu',
                #kernel_regularizer = regularization, name='conv_64x64x1_inception_in')(ir

inception_1 = build_inception_module(input_img, [64,96,128,16,32,32], 1, False, use_norm)

inception_2 = build_inception_module(inception_1, [128,128,192,32,96,64], 2, False, use_norm)

inception_3 = build_inception_module(inception_2, [192,96,208,16,48,64], 3, False, use_norm)

inception_4 = build_inception_module(inception_3, [160, 112, 224, 24, 64, 64], 4, False, use_norm)
#tower_3 = build_simple_tower(inception_2, 144, (3,3), 3, False, use_norm)
#tower_4 = build_simple_tower_subsample(tower_3, 144, (3,3), 4, False, use_norm)

#tower_5 = build_simple_tower(tower_4, 288, (3,3), 5, False, use_norm)
#tower_6 = build_simple_tower_subsample(tower_5, 288, (3,3), 6, False, use_norm)

#model top

flat_pool = AveragePooling2D(pool_size=(2, 2), padding='valid')(inception_4)

flat = Flatten()(flat_pool)

```

```

dense_5 = build_dense(flat, 128, 1, True, use_norm)

dense_6 = build_dense(dense_5, 64, 2, True, use_norm)

out = Dense(10, activation='softmax')(dense_6)

model = Model(inputs = input_img, outputs = out)

#-----

model.compile(loss='binary_crossentropy',
              optimizer=Adam(lr=0.001),
              metrics=['accuracy'])

model.summary()

model_json = model.to_json()
with open("./models/"+str(i)+".json", "w") as json_file:
    json_file.write(model_json)

print("Saved model to" + "../models/"+str(i)+".json")

```



WARNING:tensorflow:From /tensorflow-1.15.2/python3.6/tensorflow_core/python/ops/resol
Instructions for updating:

If using Keras pass *_constraint arguments to layers.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:75: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:77: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:79: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:81: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:83: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:85: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:87: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:75: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:77: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:79: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:81: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:83: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:85: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:87: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:75: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:77: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:79: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:81: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:83: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:85: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:87: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:75: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:77: UserWarning: Update
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:79: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:81: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:83: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:85: UserWarning: Update
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:87: UserWarning: Update
Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
=====			
input (InputLayer)	(None, 32, 32, 3)	0	
inception_1_3x3_reduce (Conv2D)	(None, 32, 32, 96)	384	input[0][0]
inception_1_5x5_reduce (Conv2D)	(None, 32, 32, 16)	64	input[0][0]
inception_1_pool (MaxPooling2D)	(None, 32, 32, 3)	0	input[0][0]
inception_1_1x1 (Conv2D)	(None, 32, 32, 64)	256	input[0][0]
inception_1_3x3 (Conv2D)	(None, 32, 32, 128)	110720	inception_1_3x3_rec
inception_1_5x5 (Conv2D)	(None, 32, 32, 32)	12832	inception_1_5x5_rec
inception_1_pool_proj (Conv2D)	(None, 32, 32, 32)	128	inception_1_pool[0]
inception_1_output (Concatenat	(None, 32, 32, 256)	0	inception_1_1x1[0][inception_1_3x3[0][inception_1_5x5[0][inception_1_pool_pr
inception_1_output_norm (Batch	(None, 32, 32, 256)	1024	inception_1_output[

inception_1_2x2subsample	(MaxPo (None, 16, 16, 256)	0	inception_1_/output_
inception_2_/3x3_reduce	(Conv2D (None, 16, 16, 128)	32896	inception_1_2x2subsa
inception_2_/5x5_reduce	(Conv2D (None, 16, 16, 32)	8224	inception_1_2x2subsa
inception_2_/pool	(MaxPooling2D (None, 16, 16, 256)	0	inception_1_2x2subsa
inception_2_/1x1	(Conv2D) (None, 16, 16, 128)	32896	inception_1_2x2subsa
inception_2_/3x3	(Conv2D) (None, 16, 16, 192)	221376	inception_2_/3x3_rec
inception_2_/5x5	(Conv2D) (None, 16, 16, 96)	76896	inception_2_/5x5_rec
inception_2_/pool_proj	(Conv2D) (None, 16, 16, 64)	16448	inception_2_/pool[0]
inception_2_/output	(Concatenat (None, 16, 16, 480)	0	inception_2_/1x1[0][inception_2_/3x3[0][inception_2_/5x5[0][inception_2_/pool_pr
inception_2_/output_norm	(Batch (None, 16, 16, 480)	1920	inception_2_/output[
inception_2_2x2subsample	(MaxPo (None, 8, 8, 480)	0	inception_2_/output_
inception_3_/3x3_reduce	(Conv2D (None, 8, 8, 96)	46176	inception_2_2x2subsa
inception_3_/5x5_reduce	(Conv2D (None, 8, 8, 16)	7696	inception_2_2x2subsa
inception_3_/pool	(MaxPooling2D (None, 8, 8, 480)	0	inception_2_2x2subsa
inception_3_/1x1	(Conv2D) (None, 8, 8, 192)	92352	inception_2_2x2subsa
inception_3_/3x3	(Conv2D) (None, 8, 8, 208)	179920	inception_3_/3x3_rec
inception_3_/5x5	(Conv2D) (None, 8, 8, 48)	19248	inception_3_/5x5_rec
inception_3_/pool_proj	(Conv2D) (None, 8, 8, 64)	30784	inception_3_/pool[0]
inception_3_/output	(Concatenat (None, 8, 8, 512)	0	inception_3_/1x1[0][inception_3_/3x3[0][inception_3_/5x5[0][inception_3_/pool_pr
inception_3_/output_norm	(Batch (None, 8, 8, 512)	2048	inception_3_/output[
inception_3_2x2subsample	(MaxPo (None, 4, 4, 512)	0	inception_3_/output_
inception_4_/3x3_reduce	(Conv2D (None, 4, 4, 112)	57456	inception_3_2x2subsa
inception_4_/5x5_reduce	(Conv2D (None, 4, 4, 24)	12312	inception_3_2x2subsa
inception_4_/pool	(MaxPooling2D (None, 4, 4, 512)	0	inception_3_2x2subsa
inception_4_/1x1	(Conv2D) (None, 4, 4, 160)	82080	inception_3_2x2subsa
inception_4_/3x3	(Conv2D) (None, 4, 4, 224)	226016	inception_4_/3x3_rec
inception_4_/5x5	(Conv2D) (None, 4, 4, 64)	38464	inception_4_/5x5_rec
inception_4_/pool_proj	(Conv2D) (None, 4, 4, 64)	32832	inception_4_/pool[0]

```
inception_v1/pool_proj (conv2d/ (None, 1, 1, 3), 0.001, 0.001, 0.001) 0.001
inception_v1/output/Concatenat (None, 4, 4, 512) 0.001
inception_v1/output/Concatenat (None, 4, 4, 512) 0.001
```

```
import tensorflow as tf

with tf.device('/gpu:0'):
    model.fit(x_train, y_train_cat, batch_size=128, epochs=100, validation_split=0.2, verbose=1)

result = model.evaluate(x_test, y_test_cat)

print("Accuracy on test set: ", result[1]*100, "%")
```




```
WARNING:tensorflow:From /tensorflow-1.15.2/python3.6/tensorflow_core/python/ops/math_
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Train on 40000 samples, validate on 10000 samples
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/callbacks/tensor

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/callbacks/tensor

Epoch 1/100
40000/40000 [=====] - 82s 2ms/step - loss: 0.9786 - accuracy

Epoch 00001: val_loss improved from inf to 0.54750, saving model to ./models/cifar10-
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/callbacks/tensor

Epoch 2/100
40000/40000 [=====] - 72s 2ms/step - loss: 0.3377 - accuracy

Epoch 00002: val_loss improved from 0.54750 to 0.33929, saving model to ./models/cifar10-
Epoch 3/100
40000/40000 [=====] - 72s 2ms/step - loss: 0.2591 - accuracy

Epoch 00003: val_loss improved from 0.33929 to 0.31674, saving model to ./models/cifar10-
Epoch 4/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.2296 - accuracy

Epoch 00004: val_loss improved from 0.31674 to 0.24687, saving model to ./models/cifar10-
Epoch 5/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.2144 - accuracy

Epoch 00005: val_loss did not improve from 0.24687
Epoch 6/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.2049 - accuracy

Epoch 00006: val_loss improved from 0.24687 to 0.22281, saving model to ./models/cifar10-
Epoch 7/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.1969 - accuracy

Epoch 00007: val_loss did not improve from 0.22281
Epoch 8/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.1887 - accuracy

Epoch 00008: val_loss did not improve from 0.22281
Epoch 9/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.1846 - accuracy

Epoch 00009: val_loss improved from 0.22281 to 0.18743, saving model to ./models/cifar10-
Epoch 10/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.1789 - accuracy

Epoch 00010: val_loss did not improve from 0.18743
Epoch 11/100
40000/40000 [=====] - 72s 2ms/step - loss: 0.1765 - accuracy

Epoch 00011: val_loss did not improve from 0.18743
Epoch 12/100
40000/40000 [=====] - 72s 2ms/step - loss: 0.1687 - accuracy

Epoch 00012: val_loss did not improve from 0.18743
Epoch 13/100
```

```

40000/40000 [=====] - 72s 2ms/step - loss: 0.1663 - accuracy

Epoch 00013: val_loss did not improve from 0.18743
Epoch 14/100
40000/40000 [=====] - 72s 2ms/step - loss: 0.1645 - accuracy

Epoch 00014: val_loss did not improve from 0.18743
Epoch 15/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.1309 - accuracy

Epoch 00015: val_loss improved from 0.18743 to 0.14383, saving model to ./models/cifar10-nrcrt7-04:40PM_May-08-2020.hdf5
Epoch 16/100
40000/40000 [=====] - 72s 2ms/step - loss: 0.1107 - accuracy

Epoch 00016: val_loss improved from 0.14383 to 0.13995, saving model to ./models/cifar10-nrcrt7-04:40PM_May-08-2020.hdf5
Epoch 17/100
40000/40000 [=====] - 72s 2ms/step - loss: 0.0990 - accuracy

Epoch 00017: val_loss improved from 0.13995 to 0.13503, saving model to ./models/cifar10-nrcrt7-04:40PM_May-08-2020.hdf5
Epoch 18/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.0906 - accuracy

Epoch 00018: val_loss improved from 0.13503 to 0.13293, saving model to ./models/cifar10-nrcrt7-04:40PM_May-08-2020.hdf5
Epoch 19/100
40000/40000 [=====] - 71s 2ms/step - loss: 0.0834 - accuracy

```

```

predict = model.predict(x_test)
m = max(predict[7])
index = [i for i,j in enumerate(predict[7]) if j == m]
print("The value of the prediction of test sample with index 7 is :")
print(index)

```

```

☞ The value of the prediction of test sample with index 7 is :
[6]

```

```


print("The actual class of test sample with index 7 is: ")
plt.imshow(x_test[7,:,:,:], cmap='gray')
plt.show()

```

```

☞ The actual class of test sample with index 7 is:

```



```

40000/40000 [=====] - 72s 2ms/step - loss: 0.0507 - accuracy

```

```

model.load_weights('./models/cifar10-nrcrt7-04:40PM_May-08-2020.hdf5')

```

```

result = model.evaluate(x_test, y_test_cat)

```

```
print(result)
```

```
☞ 10000/10000 [=====] - 8s 794us/step  
[0.13478883649110793, 0.9691197276115417]
```

```
Epoch 00029: early stopping
```

```
10000/10000 [=====] - 8s 815us/step
```

```
Accuracy on test set: 96.87296748161316 %
```